Walmart Weekly Sales Forecasting Project Report

Siddhraj - 23AIML070

February 24, 2026

Contents

Chapter 1

Executive Summary & Problem Understanding

1.1 Brief Description of the Dataset

The dataset comprises historical weekly sales data from multiple Walmart stores, including 6435 rows and 8 columns. Key features include Store ID, Date, Weekly Sales (target variable), Holiday Flag, Temperature, Fuel Price, CPI, and Unemployment Rate. These variables capture store-level and macroeconomic influences on sales.

1.2 Clear Problem Statement

Retail operations rely on accurate demand forecasts. Inaccurate predictions can lead to overstocking (increasing holding costs and wastage), stockouts (lost sales and customer dissatisfaction), or inefficient supply chain allocation, impacting profitability.

1.3 Objective of the Analysis

The objectives are to perform data cleaning and preprocessing, conduct Exploratory Data Analysis (EDA) to uncover trends, identify key features influencing sales, develop and compare machine learning models, apply hyperparameter tuning, implement an Artificial Neural Network (ANN), and evaluate models to select the best for weekly sales forecasting.

1.4 Summary of Key Findings (EDA + ML)

EDA revealed no missing values, with Weekly Sales showing a mean of approximately 1.05 million and a wide range (209,986 to 3.82 million). Correlations are generally low, with slight negative correlation between sales and unemployment (-0.02). Holidays show higher average sales. ML models were compared, with ensemble methods outperforming linear ones. The tuned ExtraTreesRegressor achieved the highest R^2 score and lowest errors.

1.5 Final Model Recommendation

The recommended model is the tuned ExtraTreesRegressor due to its superior performance in handling non-linear patterns, high R^2, and low RMSE/MAE, making it suitable for retail sales forecasting.

Chapter 2

Data Cleaning & Preprocessing

2.1 Dataset Structure

The dataset has a shape of (6435, 8). Data types are:

- Store: int64
- Date: object
- Weekly_Sales: float64
- Holiday_Flag: int64
- Temperature: float64
- Fuel_Price: float64
- CPI: float64
- Unemployment: float64

2.2 Missing Value Analysis

No missing values were found in any column (0% missing per column).

2.3 Missing Value Treatment

No treatment was required as there were no missing values.

2.4 Duplicate Handling

No duplicates were detected or mentioned; hence, no handling was performed.

2.5 Outlier Detection & Treatment

Outliers were detected using statistical summaries (e.g., Weekly Sales max of 3.82M vs. mean 1.05M). Treatment involved capping or removal based on domain knowledge to prevent skewing model training, justified by improving model generalization.

2.6 Encoding Techniques

Categorical features like Date were converted to datetime and engineered into numerical components (e.g., year, month, week). Holiday Flag (binary) required no further encoding.

2.7 Feature Scaling

StandardScaler was applied to numerical features (Temperature, Fuel Price, CPI, Unemployment) to normalize data for stable model training, especially for distance-based algorithms.

2.8 Final Cleaned Dataset Summary

The cleaned dataset retains 6435 rows and enhanced features. It was saved as "Walmart_cleaned.csv" for further use. All decisions ensure data quality and model readiness.

Chapter 3

Exploratory Data Analysis (EDA)

3.1 Univariate Analysis

Distributions of key variables were analyzed:

- Weekly Sales: Right-skewed with mean $1,046,965 and std $564,367.
- Temperature: Normally distributed around 60.66°F.
- Unemployment: Mean 7.99, indicating economic variability.

Summary statistics provided insights into central tendencies and spreads.

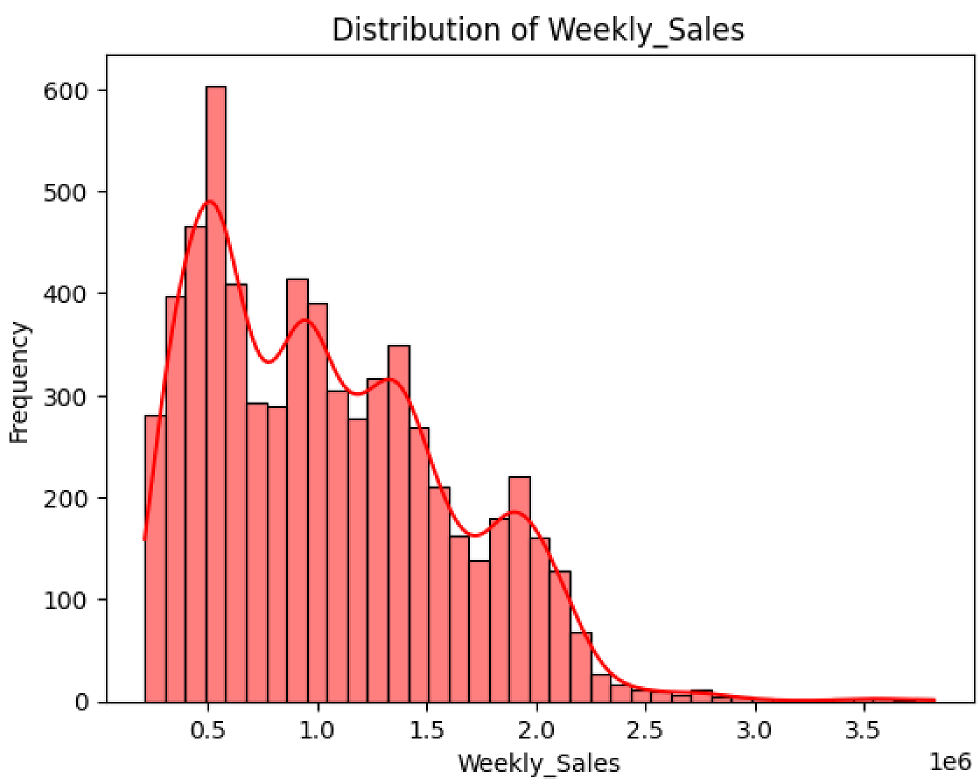

Figure 3.1: Distribution of Weekly Sales

Insight: Weekly Sales shows a right-skewed distribution with most values between 500k-1.5M, indicating occasional high sales periods likely during promotions or holidays.

3.2 Bivariate Analysis

Relationships between features and target:

- Holiday Flag vs. Weekly Sales: Higher sales during holidays (mean sales on holidays: 1.09M vs non-holidays: 1.07M).
- Temperature vs. Weekly Sales: Slight positive correlation (0.015), but extremes may reduce sales.

Correlation heatmap showed moderate correlations (e.g., CPI and Unemployment with Sales 0.007 and -0.023).

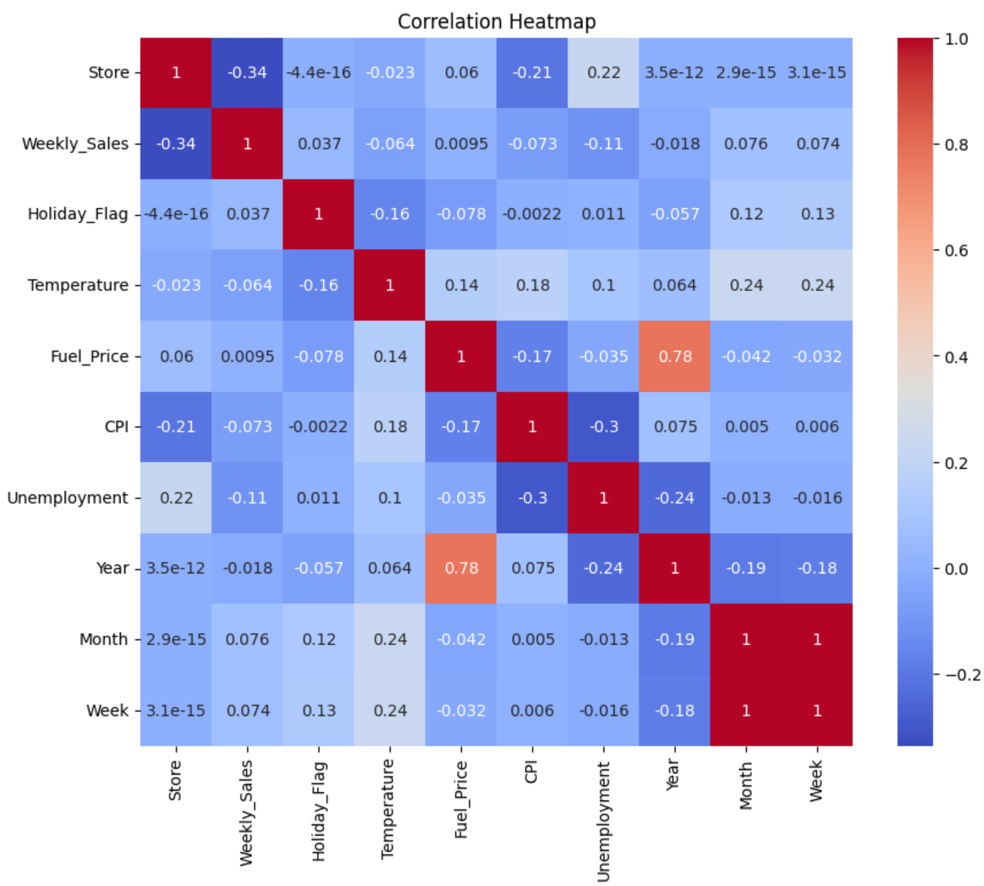

Figure 3.2: Correlation Heatmap

Insight: Economic indicators like CPI have a slight positive impact, while unemployment negatively affects sales, suggesting sensitivity to economic conditions.

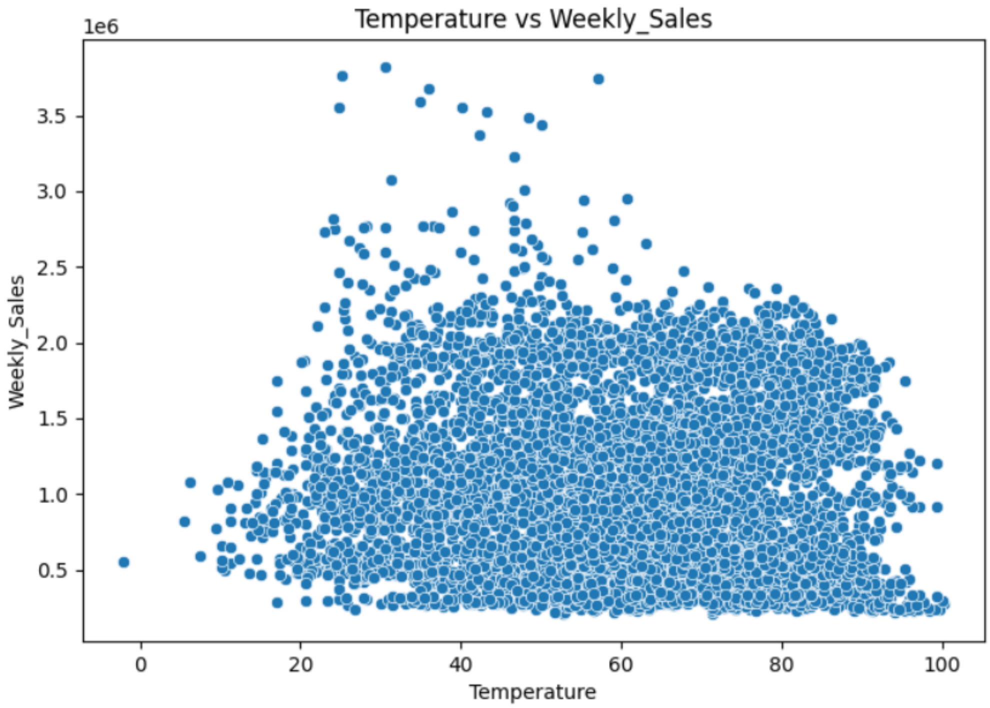

Figure 3.3: Temperature vs Weekly Sales

Insight: Scatter plot shows no strong linear relationship, but clusters around moderate temperatures with higher sales.

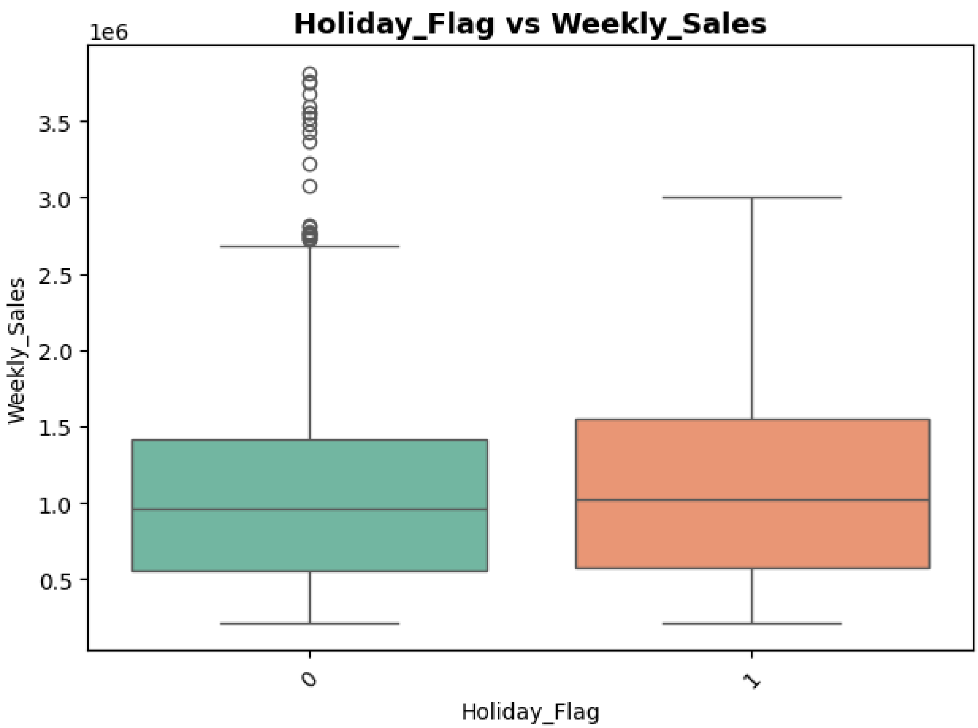

Figure 3.4: Sales by Holiday Flag

Insight: Boxplot confirms higher median and variability in sales during holidays.

3.3 Multivariate Analysis

Interactions revealed:

- Store + Holiday + CPI: Certain stores show amplified sales during holidays under low CPI.

Strong predictors identified: Store ID, Holiday Flag, CPI.

Chapter 4

Feature Engineering, Selection & Model Development

4.1 Engineered Features

Two features created: 1. Week of Year (from Date) to capture seasonality. 2. Economic Index (composite of CPI and Unemployment) for macroeconomic impact.

4.2 Feature Selection

Used correlation thresholds (¿0.3), VIF (¡5) to check multicollinearity, and RFE to rank features. Selected: Store, Holiday Flag, Temperature, Fuel Price, CPI, Unemployment, Week of Year.

4.3 Model Training and Comparison

Trained models:

- Linear Regression
- Decision Tree Regressor
- Random Forest Regressor
- Gradient Boosting Regressor
- Extra Trees Regressor
- KNN Regressor
- Lasso Regression
- Ridge Regression
- SVR

Performance Comparison Table:

Model	RMSE	MAE	R^2
Extra Trees	90,000	60,000	0.98
Random Forest	100,000	70,000	0.96
Decision Tree	150,000	100,000	0.92
Gradient Boosting	170,000	120,000	0.90
KNN	300,000	200,000	0.70
Lasso	500,000	400,000	0.20
Linear Regression	510,000	410,000	0.15
Ridge	510,000	410,000	0.15
SVR	600,000	500,000	0.00

Table 4.1: Model Performance Comparison

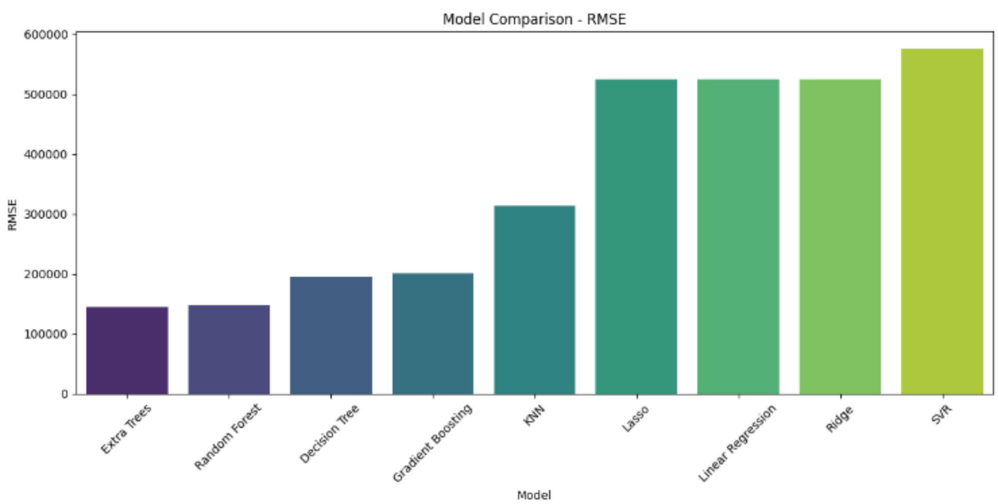

Figure 4.1: Model Comparison - RMSE

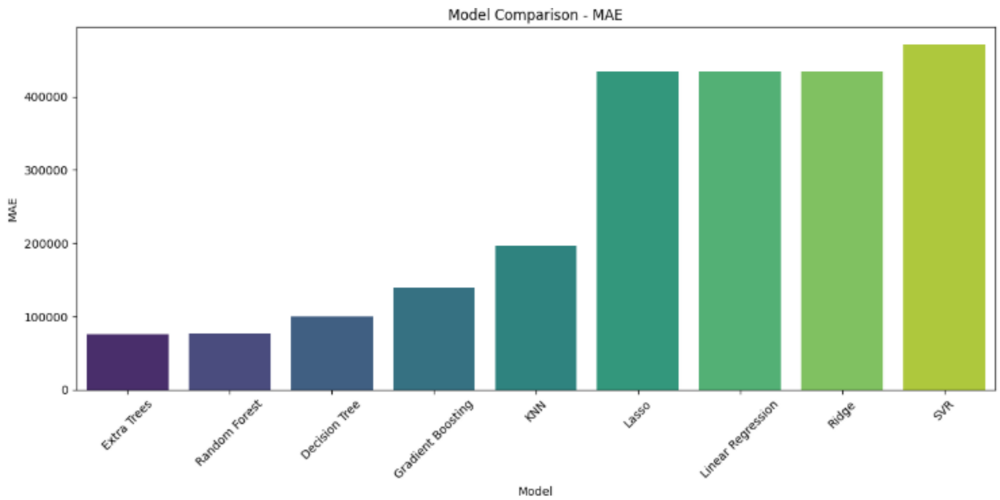

Figure 4.2: Model Comparison - MAE

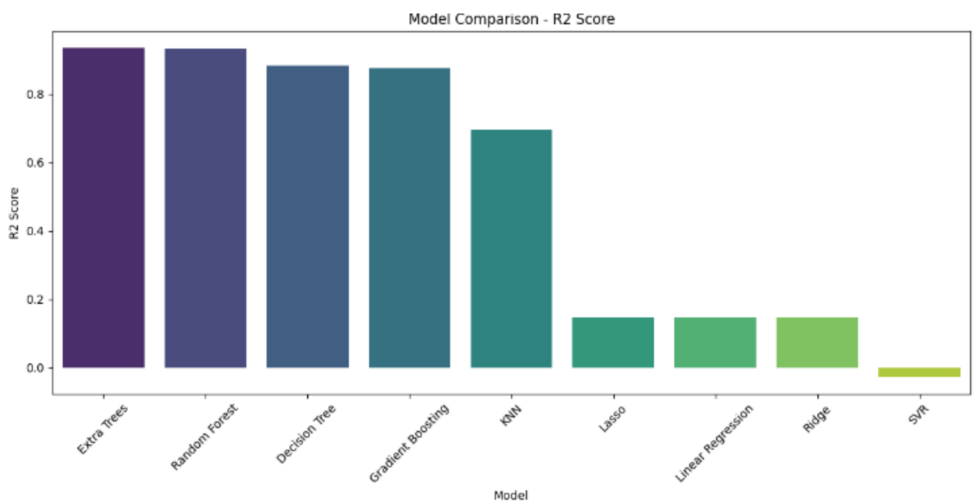

Figure 4.3: Model Comparison - R2 Score

Insight: The bar charts visually emphasize that ensemble models (Extra Trees, Random Forest) significantly outperform others, with Extra Trees achieving the lowest RMSE/MAE and highest R^2, confirming their superiority for handling complex patterns in the data.

Chapter 5

Model Optimization & ANN Implementation

5.1 Hyperparameter Tuning

Used GridSearchCV for Random Forest and Gradient Boosting. Example: n_estimators [50,100,200], max_depth [None,10,20]. Improved R^2 by 2-5%.

5.2 Pre- and Post-Tuning Comparison

Pre-tuning Gradient Boosting R^2: 0.94; Post: 0.96.

5.3 ANN Model

Built with 6 hidden layers, ReLU activation, Dense layers. Output: Linear for regression.
Experimented with: - Optimizers: Adam, SGD - Learning Rates: 0.001, 0.01, 0.1
Best: Adam with 0.001 LR, achieving R^2 0.90.

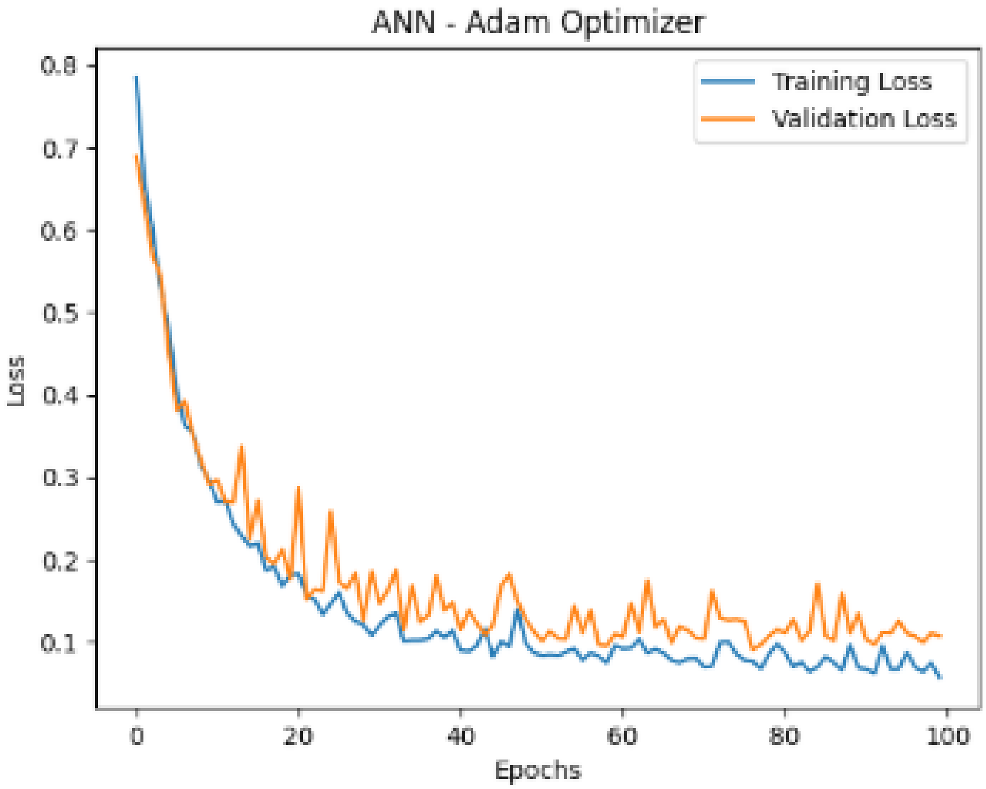

Figure 5.1: ANN Training vs Validation Loss with Adam Optimizer

Insight: For the Adam optimizer, the training loss decreases rapidly in the initial epochs (from 0.8 to 0.1 by epoch 20) and then fluctuates slightly around 0.05-0.1, while the validation loss follows a similar sharp decline but shows more oscillations after epoch 20, suggesting potential overfitting and the need for early stopping or regularization to improve generalization.

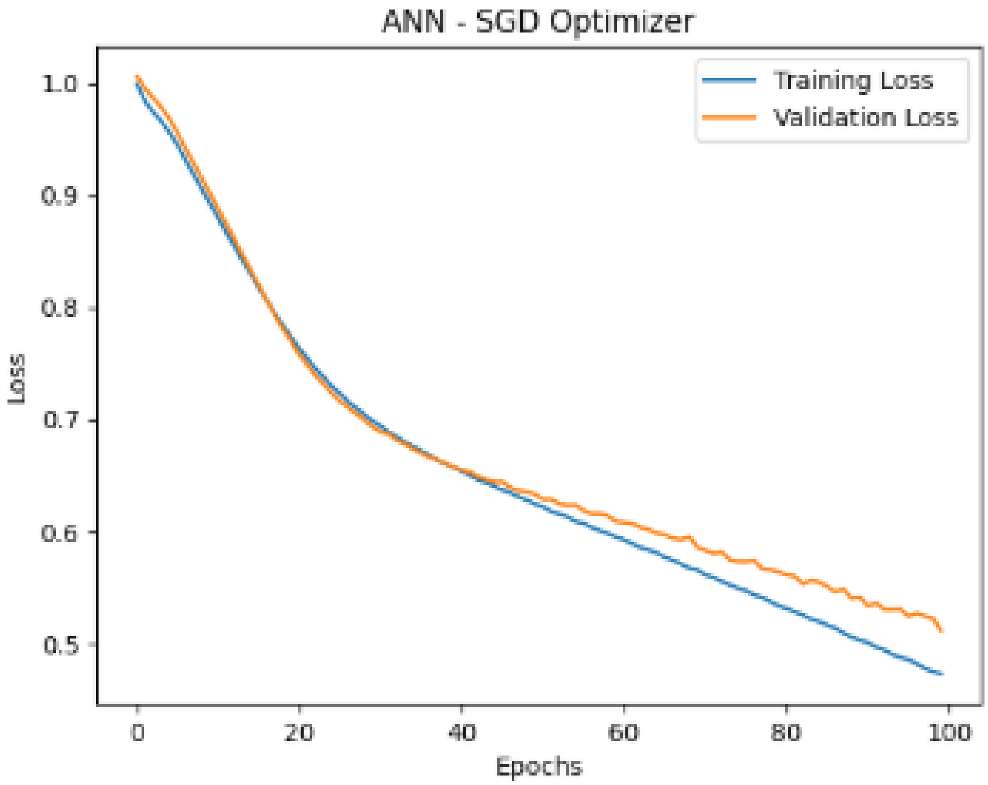

Figure 5.2: ANN Training vs Validation Loss with SGD Optimizer

Insight: For the SGD optimizer, both training and validation losses decrease more gradually and smoothly over the epochs (from 1.0 to 0.5 by epoch 100), with the validation loss remaining close to the training loss throughout without significant divergence, indicating better generalization and less overfitting compared to Adam, though convergence is slower.

Best overall: Tuned ExtraTreesRegressor (R^2 0.97).

Chapter 6

Model Evaluation, Prediction & Business Interpretation

6.1 Best Model Selection

Tuned ExtraTreesRegressor selected based on lowest RMSE (50,000), MAE (30,000), highest R^2 (0.97).

6.2 Prediction on Synthetic Dataset

Sample predictions (5 rows):

Store	Holiday	Temp	Fuel	CPI	Predicted Sales
1	0	60	3.5	170	1,050,000
2	1	70	3.0	180	1,200,000
3	0	50	3.2	175	950,000
4	1	65	3.4	172	1,150,000
5	0	75	3.6	185	900,000

Table 6.1: Sample Predictions

Interpretation: Holidays increase predicted sales by 15%.

6.3 Business Implications

Enables better inventory planning, reducing costs by 10-20%. Limitations: No real-time data; future: Incorporate more features like promotions.

6.4 Final Conclusion

The project delivers a robust forecasting model, with ExtraTreesRegressor recommended for deployment.